

Potato Disease Classification

Import all the Libraries

```
In [ ]: import tensorflow as tf
        from tensorflow.keras import models, layers
        import matplotlib.pyplot as plt
```

```
In [2]: BATCH_SIZE = 32
        IMAGE_SIZE = 256
        CHANNELS=3
        EPOCHS=30
```

Import data into tensorflow dataset object

```
In [3]: dataset = tf.keras.preprocessing.image_dataset_from_directory(
        "mrdata",
        seed=1,
        shuffle=True,
        image_size=(IMAGE_SIZE, IMAGE_SIZE),
        batch_size=BATCH_SIZE
    )
```

Found 2152 files belonging to 3 classes.

```
In [4]: class_names = dataset.class_names
        class_names
```

```
Out[4]: ['Potato__Early_blight', 'Potato__Late_blight', 'Potato__healthy']
```


```
In [5]: for image_batch, labels_batch in dataset.take(1):
        print(image_batch.shape)
        print(labels_batch.numpy())
```

(32, 256, 256, 3)

[0 0 1 1 1 0 1 1 1 0 1 1 0 0 0 1 1 1 2 0 2 1 0 1 1 0 1 0 1 0 0 0]

displaying some random pictures

```
In [6]: plt.figure(figsize=(15, 15))
        for image_batch, labels_batch in dataset.take(1):
            for i in range(6):
                ax = plt.subplot(3, 3, i + 1)
                plt.imshow(image_batch[i].numpy().astype("uint8"))
                plt.title(class_names[labels_batch[i]])
                plt.axis("off")
```



splitting data_set

```
In [7]: train_size = 0.8
len(dataset)*train_size
```

```
Out[7]: 54.400000000000006
```

```
In [8]: train_dataset = dataset.take(54)
len(train_dataset)
```

```
Out[8]: 54
```

```
In [9]: test_dataset = dataset.skip(54)
len(test_dataset)
```

```
Out[9]: 14
```

```
In [10]: val_size=0.1
len(dataset)*val_size
```

```
Out[10]: 6.800000000000001
```

```
In [11]: val_dataset = test_dataset.take(6)
len(val_dataset)
```

```
Out[11]: 6
```

```
In [12]: test_dataset = test_dataset.skip(6)
len(test_dataset)
```

Out[12]: 8

Building the Model Creating a Layer for Resizing and Normalization Before we feed our images to network, we should be resizing it to the desired size. Moreover, to improve model performance, we should normalize the image pixel value (keeping them in range 0 and 1 by dividing by 256). This should happen while training as well as inference. Hence we can add that as a layer in our Sequential Model. Building the Model Creating a Layer for Resizing and Normalization Before we feed our images to network, we should be resizing it to the desired size. Moreover, to improve model performance, we should normalize the image pixel value (keeping them in range 0 and 1 by dividing by 256). This should happen while training as well as inference. Hence we can add that as a layer in our Sequential Model.

```
In [13]: resize_and_rescale = tf.keras.Sequential([
    layers.experimental.preprocessing.Resizing(IMAGE_SIZE, IMAGE_SIZE),
    layers.experimental.preprocessing.Rescaling(1./255),
])
```

Model Architecture We use a CNN coupled with a Softmax activation in the output layer. We also add the initial layers for resizing, normalization and Data Augmentation.

```
In [14]: input_shape = (BATCH_SIZE, IMAGE_SIZE, IMAGE_SIZE, CHANNELS)
n_classes = 3

model = models.Sequential([
    resize_and_rescale,
    layers.Conv2D(32, kernel_size = (3,3), activation='relu', input_shape=input_shape),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(n_classes, activation='softmax'),
])

model.build(input_shape=input_shape)
```

```
In [15]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(32, 256, 256, 3)	0
conv2d (Conv2D)	(32, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(32, 127, 127, 32)	0
conv2d_1 (Conv2D)	(32, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(32, 62, 62, 64)	0
conv2d_2 (Conv2D)	(32, 60, 60, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(32, 30, 30, 64)	0
conv2d_3 (Conv2D)	(32, 28, 28, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(32, 14, 14, 64)	0
conv2d_4 (Conv2D)	(32, 12, 12, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(32, 6, 6, 64)	0
conv2d_5 (Conv2D)	(32, 4, 4, 64)	36928
max_pooling2d_5 (MaxPooling2D)	(32, 2, 2, 64)	0
flatten (Flatten)	(32, 256)	0
dense (Dense)	(32, 64)	16448
dense_1 (Dense)	(32, 3)	195

```

=====
Total params: 183,747
Trainable params: 183,747
Non-trainable params: 0

```

Compiling the Model We use adam Optimizer, SparseCategoricalCrossentropy for losses, accuracy as a metric

```

In [16]: model.compile(
          optimizer='adam',
          loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
          metrics=['accuracy']
        )

```

```

In [19]: history = model.fit(
          train_dataset,

```

```

    batch_size=BATCH_SIZE,
    validation_data=val_dataset,
    verbose=1,
    epochs=13,
)

```

Epoch 1/13

54/54 [=====] - 144s 3s/step - loss: 0.1134 - accuracy: 0.95
54 - val_loss: 0.0705 - val_accuracy: 0.9740

Epoch 2/13

54/54 [=====] - 146s 3s/step - loss: 0.0690 - accuracy: 0.97
69 - val_loss: 0.2391 - val_accuracy: 0.9219

Epoch 3/13

54/54 [=====] - 145s 3s/step - loss: 0.0736 - accuracy: 0.97
11 - val_loss: 0.0680 - val_accuracy: 0.9844

Epoch 4/13

54/54 [=====] - 145s 3s/step - loss: 0.0528 - accuracy: 0.98
32 - val_loss: 0.0582 - val_accuracy: 0.9844

Epoch 5/13

54/54 [=====] - 144s 3s/step - loss: 0.1112 - accuracy: 0.95
72 - val_loss: 0.0707 - val_accuracy: 0.9844

Epoch 6/13

54/54 [=====] - 145s 3s/step - loss: 0.0620 - accuracy: 0.97
63 - val_loss: 0.1083 - val_accuracy: 0.9583

Epoch 7/13

54/54 [=====] - 148s 3s/step - loss: 0.0358 - accuracy: 0.99
07 - val_loss: 0.0324 - val_accuracy: 0.9844

Epoch 8/13

54/54 [=====] - 147s 3s/step - loss: 0.0175 - accuracy: 0.99
48 - val_loss: 0.0252 - val_accuracy: 0.9844

Epoch 9/13

54/54 [=====] - 148s 3s/step - loss: 0.0891 - accuracy: 0.96
59 - val_loss: 0.1009 - val_accuracy: 0.9583

Epoch 10/13

54/54 [=====] - 143s 3s/step - loss: 0.0650 - accuracy: 0.97
22 - val_loss: 0.1069 - val_accuracy: 0.9583

Epoch 11/13

54/54 [=====] - 146s 3s/step - loss: 0.0268 - accuracy: 0.98
84 - val_loss: 0.0288 - val_accuracy: 0.9844

Epoch 12/13

54/54 [=====] - 146s 3s/step - loss: 0.0530 - accuracy: 0.98
03 - val_loss: 0.0598 - val_accuracy: 0.9635

Epoch 13/13

54/54 [=====] - 147s 3s/step - loss: 0.0473 - accuracy: 0.98
21 - val_loss: 0.0189 - val_accuracy: 0.9948

```

In [20]: scores = model.evaluate(test_dataset)
         scores

```

8/8 [=====] - 7s 625ms/step - loss: 0.0556 - accuracy: 0.982
8

```

Out[20]: [0.05564812943339348, 0.982758641242981]

```

In []: We can see above that we get 98.00% accuracy for our test dataset. This is considered

```

In [26]: import numpy as np
         for images_batch, labels_batch in test_dataset.take(1):

             first_image = images_batch[3].numpy().astype('uint8')


```

```
first_label = labels_batch[3].numpy()

print("first image to predict")
plt.imshow(first_image)
print("actual label:", class_names[first_label])

batch_prediction = model.predict(images_batch)
print("predicted label:", class_names[np.argmax(batch_prediction[3])])
```

first image to predict
actual label: Potato__Early_blight
predicted label: Potato__Early_blight



In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: